

Samuel Stevens

samuel.robert.stevens@gmail.com · <https://samuelstevens.me>

RESEARCH VISION

My PhD work is in large-scale foundation modeling: I build models and the end-to-end systems that let them scale cleanly, run efficiently, and iterate quickly. My core strengths include distributed training and data pipelines ([BioCLIP](#) and [BioCLIP 2](#)), system/model codesign, interpretability tooling (e.g., [sparse autoencoders](#)), and rigorous benchmarking/evaluation ([MMMU](#), [Mind2Web](#), [BioBench](#)). Going forward, I want to apply this stack to general-domain large models with an emphasis on scaling, efficiency, and measurement that speeds up model iteration.

My PhD work improved domain-specific modeling by treating large-scale data as the primary bottleneck for capability gains. My vision is to build principled pipelines for large-scale datasets: domain-aware collection, deduplication, quality filtering, weak supervision, and targeted curation that improves signal-to-noise without sacrificing scale. I'm especially interested in structured supervision (text, images, metadata, sequences, structures, functions) that makes representation learning more transferable. The outcome I care about is a repeatable recipe for turning better data into better models, validated with tight evaluation loops.

EDUCATION

Ohio State University Ph.D. in Computer Science & Engineering Advisor: Professor Yu Su	August 2021 – Present <i>Columbus, OH</i>
Ohio State University (3.93 GPA) Honors B.S. in Computer Science & Engineering, <i>Summa Cum Laude</i> Honors Research Distinction in Computer Science & Engineering German Minor	August 2017 – May 2021 <i>Columbus, OH</i>

SELECTED PUBLICATIONS

[Towards Open-Ended Visual Scientific Discovery with Sparse Autoencoders](#)

Samuel Stevens, Jacob Beattie, Tanya Berger-Wolf and Yu Su *arXiv Preprint*, 2025

[Interpretable and Testable Vision Features via Sparse Autoencoders](#)

Samuel Stevens, Wei-Lun Chao, Tanya Berger-Wolf and Yu Su *arXiv Preprint*, 2025

[BioCLIP 2: Emergent Properties from Scaling Hierarchical Contrastive Learning](#)

Jianyang Gu, **Samuel Stevens**, Elizabeth G Campolongo, Matthew J Thompson, Net Zhang, Jiaman Wu, Andrei Kopanev, Zheda Mai, Alexander E. White, James Balhoff, Wasila Dahdul, Daniel Rubenstein, Hilmar Lapp, Tanya Berger-Wolf, Wei-Lun Chao and Yu Su *NeurIPS*, 2025 ([Spotlight](#))

[BioCLIP: A Vision Foundation Model for the Tree of Life](#)

Samuel Stevens, Jiaman Wu, Matthew J Thompson, Elizabeth G Campolongo, Chan Hee Song, David

Edward Carlyn, Li Dong, Wasila M Dahdul, Charles Stewart, Tanya Berger-Wolf, Wei-Lun Chao and Yu Su *CVPR*, 2024 (**Best Student Paper**)

SALSA FRESCA: Angular Embeddings and Pre-Training for ML Attacks on Learning With Errors

Samuel Stevens, Emily Wenger, Cathy Yuanchen Li, Niklas Nolte, Eshika Saxena, Francois Charton and Kristin Lauter *TMLR*, 2025

Memorization for Good: Encryption with Autoregressive Language Models

Samuel Stevens and Yu Su *arXiv Preprint*, 2023

An Investigation of Language Model Interpretability via Sentence Editing

Samuel Stevens and Yu Su *EMNLP BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, 2021

RESEARCH EXPERIENCE

Zoom AI Research Scientist Intern	May 2024 – August 2024 <i>Remote</i>
• Designed and implemented abstractions for programming language models more reliably.	
Meta Research Scientist Intern	May 2023 – August 2023 <i>Seattle, WA</i>
• Implemented and validated simpler architecture for the ongoing SALSA project.	
• Continued part-time from August 2023 to December 2023.	
Self-Attention Interpretability in Sequence Classification Undergraduate Honors Research Thesis	October 2019 – April 2021 <i>Columbus, OH</i>
• Fine-tuned pre-trained language models on existing sequence classification task.	
• Analyzed trends in the attention maps of BERT-based models to quantify differences in interpretability.	
• Identified literature gap and designed experiments leading to a baseline result.	

ENGINEERING EXPERIENCE

SpaceX (Starlink) Hardware Test Associate Engineer	May 2021 – August 2021 <i>Seattle, WA</i>
• Improved test code leading to \$66K expected annual savings and 25% increase in first-pass yield.	
• Developed data visualization for micropositioner actuator repeatability during testing.	
• Developed thermal camera testing prototype, reducing test time from 30+ minutes to 30 seconds.	
Microsoft Software Engineering Intern	May 2020 – August 2020 <i>Seattle, WA (Remote)</i>
• Developed new Power BI feature to improve customers' ability to diagnose data-refresh problems.	
• Ramped up quickly in a 100 million-line C++ codebase to begin delivering value immediately.	
TicketBay	January 2018 – August 2020

Lead Developer	<i>Columbus, OH</i>
<ul style="list-style-type: none"> Developed a mobile app for students to buy and sell football tickets. Facilitated the transfer of \$165K worth of tickets between more than 7K customers. Led system architecture decisions to balance ease of development and system performance. 	
GE Aviation	May 2018 – August 2018
Digital Technology Intern	<i>Cincinnati, OH</i>
<ul style="list-style-type: none"> Built an end-to-end testing solution to automate QA testing of customer-facing web application. Integrated test solution with existing CI/CD pipeline, leading to improved product quality. 	
The Ohio State University	May 2018 – November 2019
Lead Developer	<i>Columbus, OH</i>
<ul style="list-style-type: none"> Developed a webapp to distribute rich media content to support a professor's research. Managed 2 student developers in multiple time zones with agile methodology. 	
GE Aviation	May 2017 – August 2017
INTERAlliance Intern	<i>Cincinnati, OH</i>
<ul style="list-style-type: none"> In-sourced customer-facing search, targeting \$300K in savings and 40K customers. Led MongoDB integration, giving customers access to previously unsearchable documents. 	

AWARDS AND HONORS

CVPR Best Student Paper	2024
<ul style="list-style-type: none"> BioCLIP: A Vision Foundation Model for the Tree of Life 	
1st Alexa Prize Taskbot Challenge: 3rd Place	2022
<ul style="list-style-type: none"> Developed internal dashboard to support bug-fixes in near-real time. Designed and implemented automated test suite using fuzzy matching to support dynamic chatbot responses. 	

ALL PUBLICATIONS

2025

Towards Open-Ended Visual Scientific Discovery with Sparse Autoencoders
Samuel Stevens, Jacob Beattie, Tanya Berger-Wolf and Yu Su *arXiv Preprint*, 2025

BioCLIP 2: Emergent Properties from Scaling Hierarchical Contrastive Learning
Jianyang Gu, **Samuel Stevens**, Elizabeth G Campolongo, Matthew J Thompson, Net Zhang, Jiaman Wu, Andrei Kopanev, Zheda Mai, Alexander E. White, James Balhoff, Wasila Dahdul, Daniel Rubenstein, Hilmar Lapp, Tanya Berger-Wolf, Wei-Lun Chao and Yu Su *NeurIPS*, 2025 (**Spotlight**)

Interpretable and Testable Vision Features via Sparse Autoencoders
Samuel Stevens, Wei-Lun Chao, Tanya Berger-Wolf and Yu Su *arXiv Preprint*, 2025

SALSA FRESCA: Angular Embeddings and Pre-Training for ML Attacks on Learning With Errors

Samuel Stevens, Emily Wenger, Cathy Yuanchen Li, Niklas Nolte, Eshika Saxena, Francois Charton and Kristin Lauter *TMLR*, 2025

BioBENCH: A Blueprint to Move Beyond ImageNet for Scientific ML Benchmarks

Samuel Stevens *NeurIPS Imageomics Workshop*, 2025

Mind the (Data) Gap: Evaluating Vision Systems in Small Data Applications

Samuel Stevens, S M Rayeed and Jenna Kline *NeurIPS Imageomics Workshop*, 2025

2024

BioCLIP: A Vision Foundation Model for the Tree of Life

Samuel Stevens, Jiaman Wu, Matthew J Thompson, Elizabeth G Campolongo, Chan Hee Song, David Edward Carlyn, Li Dong, Wasila M Dahdul, Charles Stewart, Tanya Berger-Wolf, Wei-Lun Chao and Yu Su *CVPR*, 2024 (**Best Student Paper**)

MMMU: A Massive Multi-Discipline Multimodal Understanding and Reasoning Benchmark for Expert AGI

Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruqi Liu, Ge Zhang, **Samuel Stevens**, Dongfu Jiang, Weiming Ren, Yuxuan Sun, Cong Wei, Botao Yu, Ruibin Yuan, Renliang Sun, Ming Yin, Boyuan Zheng, Zhenzhu Yang, Yibo Liu, Wenhao Huang, Huan Sun, Yu Su and Wenhui Chen *CVPR*, 2024 (**Best Paper Finalist**)

A Simple Interpretable Transformer for Fine-Grained Image Classification and Analysis

Dipanjyoti Paul, Arpita Chowdhury, Xinqi Xiong, Feng-Ju Chang, David Carlyn, **Samuel Stevens**, Kaiya Provost, Anuj Karpatne, Bryan Carstens, Daniel Rubenstein, Charles Stewart, Tanya Berger-Wolf, Yu Su and Wei-Lun Chao *ICLR*, 2024

The Cool and the Cruel: Separating Hard Parts of LWE Secrets

Niklas Nolte, Mohamed Malhou, Emily Wenger, **Samuel Stevens**, Cathy Li, Francois Charton and Kristin Lauter *AFRICACRYPT*, 2024

KABR: In-Situ Dataset for Kenyan Animal Behavior Recognition from Drone Videos

Maksim Kholiavchenko, Jenna Kline, Michelle Ramirez, **Samuel Stevens**, Alec Sheets, Reshma Babu, Namrata Banerji, Elizabeth Campolongo, Matthew Thompson, Nina Van Tiel, Jackson Miliko, Eduardo Bessa, Isla Duporge, Tanya Berger-Wolf, Daniel Rubenstein and Charles Stewart *WACV*, 2024

Optimizing Image Capture for Computer Vision-Powered Taxonomic Identification

Alyson East, Elizabeth G. Campolongo, Luke Meyers, **Samuel Stevens**, Hilmar Lapp, Yasin Bakis, Henry Bart, Wasila Dahdul, Anuj Karpatne, Tanya Berger-Wolf and John Bradley *Methods in Ecology and Evolution*, 2024

A Framework for Autonomic Computing for In Situ Imageomics

Jenna Kline, Christopher Stewart, Tanya Berger-Wolf, **Samuel Stevens** and Charles Stewart *IEEE ACSOS*, 2024

2023

Mind2Web: Towards a Generalist Agent for the Web

Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, **Samuel Stevens**, Boshi Wang, Huan Sun and Yu Su
NeurIPS, 2023 (Spotlight)

Memorization for Good: Encryption with Autoregressive Language Models

Samuel Stevens and Yu Su *arXiv Preprint*, 2023

Roll Up Your Sleeves: Working with a Collaborative and Engaging Task-Oriented Dialogue System

Lingbo Mo, Shijie Chen, Ziru Chen, Xiang Deng, Ashley Lewis, Sunit Singh, **Samuel Stevens**, Chang-You Tai, Zhen Wang, Xiang Yue, Tianshu Zhang, Yu Su and Huan Sun *SIGDIAL*, 2023

SALSABOT: Towards a Robust and Generalizable Embodied Agent

Chan Hee Song, Jiaman Wu, Ju-Seung Byeon, Zexin Xu, Vardaan Pahuja, Goonmeet Bajaj, **Samuel Stevens**, Ziru Chen and Yu Su *CVPR Embodied AI Workshop*, 2023

2022

arXivEdits: Understanding the Human Revision Process in Scientific Writing

Chao Jiang, Wei Xu and **Samuel Stevens** *EMNLP*, 2022

Bootstrapping a User-Centered Task-Oriented Dialogue System

Shijie Chen, Ziru Chen, Xiang Deng, Ashley Lewis, Lingbo Mo, **Samuel Stevens**, Zhen Wang, Xiang Yue, Tianshu Zhang, Yu Su and Huan Sun *Alexa Prize TaskBot Challenge Proceedings*, 2022

2021

An Investigation of Language Model Interpretability via Sentence Editing

Samuel Stevens and Yu Su *EMNLP BlackboxNLP Workshop*, 2021